

# SQL STORED ROUTINES

CS121: Introduction to Relational Database Systems  
Fall 2016 – Lecture 9

# SQL Functions

2

- SQL queries can use sophisticated math operations and functions
  - ▣ Can compute simple functions, aggregates
  - ▣ Can compute and filter results
- Sometimes, apps require specialized computations
  - ▣ Would like to use these in SQL queries, too
- SQL provides a mechanism for defining functions
  - ▣ Called User-Defined Functions (UDFs)

# SQL Functions (2)

3

- Can be defined in a procedural SQL language, or in an external language
  - ▣ SQL:1999, SQL:2003 both specify a language for declaring functions and procedures
- Different vendors provide their own languages
  - ▣ Oracle: PL/SQL
  - ▣ Microsoft: Transact-SQL (T-SQL)
  - ▣ PostgreSQL: PL/pgSQL
  - ▣ MySQL: stored procedure support strives to follow specifications (and mostly does)
  - ▣ Some also support external languages: Java, C, C#, etc.
- As usual, lots of variation in features and syntax

# Example SQL Function

4

- A SQL function to count how many bank accounts a particular customer has:

```
CREATE FUNCTION account_count(  
    customer_name VARCHAR(20)  
) RETURNS INTEGER  
BEGIN  
    DECLARE a_count INTEGER;  
  
    SELECT COUNT(*) INTO a_count FROM depositor AS d  
    WHERE d.customer_name = customer_name;  
  
    RETURN a_count;  
END
```

- ▣ Function can take arguments and return values
- ▣ Can use SQL statements and other operations in body

# Example SQL Function (2)

5

- Can use our function for individual accounts:

```
SELECT account_count('Johnson');
```

- Can include in computed results:

```
SELECT customer_name,  
       account_count(customer_name) AS accts  
FROM customer;
```

- Can include in **WHERE** clause:

```
SELECT customer_name FROM customer  
WHERE account_count(customer_name) > 1;
```

# Arguments and Return-Values

6

- ❑ Functions can take any number of arguments (even 0)
- ❑ Functions *must* return a value
  - ▣ Specify type of value in **RETURNS** clause
- ❑ From our example:

```
CREATE FUNCTION account_count(  
    customer_name VARCHAR(20)  
) RETURNS INTEGER
```

  - ▣ One argument named **customer\_name**, type is **VARCHAR(20)**
  - ▣ Returns some **INTEGER** value

# Table Functions

7

- SQL:2003 spec. includes table functions
  - ▣ Return a whole table as their result
  - ▣ Can be used in **FROM** clause
- A generalization of views
  - ▣ Can be considered to be parameterized views
  - ▣ Call function with specific arguments
  - ▣ Result is a relation based on those arguments
- Although SQL:2003 not broadly supported yet, most DBMSes provide a feature like this
  - ▣ *...in various ways, of course...*

# Function Bodies and Variables

8

- Blocks of procedural SQL commands are enclosed with **BEGIN** and **END**
  - ▣ Defines a compound statement
  - ▣ Can have nested **BEGIN ... END** blocks
- Variables are specified with **DECLARE** statement
  - ▣ Must appear at start of a block
  - ▣ Initial value is **NULL**
  - ▣ Can initialize to some other value with **DEFAULT** syntax
  - ▣ Scope of a variable is within its block
  - ▣ Variables in inner blocks can shadow variables in outer blocks



# Example Blocks and Variables

9

- Our `account_count` function's body:

```
BEGIN
```

```
    DECLARE a_count INTEGER;
```

```
    SELECT COUNT(*) INTO a_count FROM depositor AS d
    WHERE d.customer_name = customer_name;
```

```
    RETURN a_count;
```

```
END
```

- A simple integer variable with initial value:

```
BEGIN
```

```
    DECLARE result INTEGER DEFAULT 0;
```

```
    ...
```

```
END
```

# Assigning To Variables

10

- Can use **SELECT ... INTO** syntax

- ▣ For assigning the result of a query into a variable

```
SELECT COUNT(*) INTO a_count
```

```
FROM depositor AS d
```

```
WHERE d.customer_name = customer_name;
```

- ▣ Query must produce a single row

**Note: SELECT INTO** sometimes has multiple meanings!

This form is specific to the body of stored routines.

- e.g. frequently used to create a temp table from a **SELECT**

- Can also use **SET** syntax

- ▣ For assigning result of a math expression to a variable

```
SET result = n * (n + 1) / 2;
```

# Assigning Multiple Variables

11

- Can assign to multiple variables using **SELECT INTO** syntax
- Example: Want both the number of accounts and the total balance

```
DECLARE a_count INTEGER;  
DECLARE total_balance NUMERIC(12,2);  
  
SELECT COUNT(*), SUM(balance)  
INTO a_count, total_balance  
FROM depositor AS d NATURAL JOIN account  
WHERE d.customer_name = customer_name;
```

# Another Example

12

- Simple function to compute sum of 1..N

```
CREATE FUNCTION sum_n(n INTEGER) RETURNS INTEGER
BEGIN
    DECLARE result INTEGER DEFAULT 0;
    SET result = n * (n + 1) / 2;
    RETURN result;
END
```

- Lots of extra work in that! To simplify:

```
CREATE FUNCTION sum_n(n INTEGER) RETURNS INTEGER
BEGIN
    RETURN n * (n + 1) / 2;
END
```

# Dropping Functions

13

- Can't simply overwrite functions in the database
  - ▣ Same as tables, views, etc.

- First, drop old version of function:

```
DROP FUNCTION sum_n;
```

- Then create new version of function:

```
CREATE FUNCTION sum_n(n INTEGER)
RETURNS INTEGER
BEGIN
    RETURN n * (n + 1) / 2;
END
```

# SQL Procedures

14

- Functions have specific limitations
  - ▣ Must return a value
  - ▣ All arguments are input-only
  - ▣ Typically cannot affect current transaction status (i.e. function cannot commit, rollback, etc.)
  - ▣ Usually not allowed to modify tables, except in particular circumstances
- Stored procedures are more general constructs without these limitations
  - ▣ Generally can't be used in same places as functions
  - ▣ e.g. can't use in **SELECT** clause
  - ▣ Procedures don't return a value like functions do

# Example Procedure

15

- Write a procedure that returns both the number of accounts a customer has, and their total balance

- ▣ Results are passed back using out-parameters

```
CREATE PROCEDURE account_summary(  
    IN customer_name VARCHAR(20),  
    OUT a_count INTEGER,  
    OUT total_balance NUMERIC(12,2)  
)  
BEGIN  
    SELECT COUNT(*), SUM(balance)  
    INTO a_count, total_balance  
    FROM depositor AS d NATURAL JOIN account  
    WHERE d.customer_name = customer_name;  
END
```

- Default parameter type is **IN**

# Calling a Procedure

16

- Use the **CALL** statement to invoke a procedure

```
CALL account_summary(...);
```

- To use this procedure, must also have variables to receive the values

- MySQL SQL syntax:

```
CALL account_summary('Johnson',  
                      @j_count, @j_total);
```

```
SELECT @j_count, @j_total;
```

- ▣ **@var** declares a temporary session variable

+-----+	+-----+	
@j_cnt	@j_tot	
+-----+	+-----+	
2	1400.00	
+-----+	+-----+	



# Conditional Operations

17

- SQL provides an if-then-else construct

```
IF cond1 THEN command1  
ELSEIF cond2 THEN command2  
ELSE command3  
END IF
```

- ▣ Branches can also specify compound statements instead of single statements
  - Enclose compound statements with **BEGIN** and **END**
- ▣ Can leave out **ELSEIF** and/or **ELSE** clauses, as usual

# Looping Constructs

18

- SQL also provides looping constructs

- **WHILE** loop:

```
DECLARE n INTEGER DEFAULT 0;  
WHILE n < 10 DO  
    SET n = n + 1;  
END WHILE;
```

- **REPEAT** loop:

```
REPEAT  
    SET n = n - 1;  
UNTIL n = 0  
END REPEAT;
```

# Iteration Over Query Results

19

- Sometimes need to issue a query, then iterate over each row in result
  - ▣ Perform more sophisticated operations than a simple SQL query can perform
- Examples:
  - ▣ Many kinds of values that standard OLTP databases can't compute quickly!
  - ▣ Assign a rank to a collection of rows:
    - Can compare each row to all other rows, typically with a cross-join
    - Or, sort rows then iterate over results, assigning rank values
  - ▣ Given web logs containing individual HTTP request records:
    - Compute each client's "visit length," from requests that are within 20 minutes of some other request from the same client

# Cursors

20

- Need to issue a query to fetch specific results
- Then, need to iterate through each row in the result
  - ▣ Operate on each row's values individually
- A cursor is an iterator over rows in a result set
  - ▣ Cursor refers to one row in query results
  - ▣ Can access row's values through the cursor
  - ▣ Can move cursor forward through results
- Cursors can provide different features
  - ▣ Read-only vs. read-write
  - ▣ Forward-only vs. bidirectional
  - ▣ Static vs. dynamic (when concurrent changes occur)

# Cursor Notes

21

- Cursors can be expensive
- Can the operation use a normal SQL query instead?
  - ▣ (Usually, the answer is yes...)
  - ▣ Cursors let you do what databases do, but slower
- Cursors might also hold system resources until they are finished
  - ▣ e.g. DB might store query results in a temporary table, to provide a read-only, static view of query result
- Syntax varies widely across DBMSes
- Most external DB connectivity APIs provide cursor capabilities

# Stored Routines and Cursors

22

- Can use cursors inside stored procedures and UDFs
- Syntax from the book:

```
DECLARE n INTEGER DEFAULT 0;  
FOR r AS SELECT balance FROM account  
        WHERE branch_name='Perryridge'  
DO  
    SET n = n + r.balance;  
END FOR
```

- ▣ Iterates over account balances from Perryridge branch, summing balances
- ▣ `r` is implicitly a cursor
  - `FOR` construct automatically moves the cursor forward
- ▣ (Could compute this with a simple SQL query, too...)

# MySQL Cursor Syntax

23

- Must *explicitly* declare cursor variable  
`DECLARE cur CURSOR FOR  
SELECT ... ;`
- Open cursor to use query results:  
`OPEN cur;`
- Fetch values from cursor into variables  
`FETCH cur INTO var1, var2, ... ;`
  - ▣ Next row is fetched, and values are stored into specified variables
  - ▣ Must specify the same number of variables as columns in the result
  - ▣ A specific error condition is flagged to indicate end of results
- Close cursor at end of operation  
`CLOSE cur;`
  - ▣ Also happens automatically at end of enclosing block

# Handling Errors

24

- Many situations where errors can occur in stored procedures
  - ▣ Called conditions
  - ▣ Includes errors, warnings, other signals
  - ▣ Can also include user-defined conditions
- Handlers can be defined for conditions
- When a condition is signaled, its handler is invoked
  - ▣ Handler can specify whether to continue running the procedure, or whether to exit procedure instead



# Conditions

25

- Predefined conditions:

- ▣ **NOT FOUND**

- Query fetched no results, or command processed no results

- ▣ **SQLWARNING**

- Non-fatal SQL problem occurred

- ▣ **SQLException**

- Serious SQL error occurred

# Conditions (2)

26

- Can also define application-specific conditions
  - ▣ Examples:
    - “Account overdraft!”
    - “Inventory of item hit zero.”
- Syntax for declaring conditions:

```
DECLARE acct_overdraft CONDITION
DECLARE zero_inventory CONDITION
```
- Not every DBMS supports generic conditions
  - ▣ e.g. MySQL supports assigning names to existing SQL error codes, but not creating new conditions

# Handlers

27

- Can declare handlers for specific conditions
- Handler specifies statements to execute
- Handler also specifies what should happen next:
  - ▣ Continue running the procedure where it left off
  - ▣ Exit the stored procedure completely
- Syntax:
  - ▣ A continue-handler:  
`DECLARE CONTINUE HANDLER FOR condition statement`
  - ▣ An exit-handler:  
`DECLARE EXIT HANDLER FOR condition statement`
  - ▣ Can also specify a statement-block instead of an individual statement

# Handlers (2)

28

- Handlers can do very simple things
  - ▣ e.g. set a flag to indicate some situation
- Can also do very complicated things
  - ▣ e.g. insert rows into other tables to log failure situations
  - ▣ e.g. properly handle an overdrawn account

# Total Account Balance – MySQL

29

- Declared as a function – returns a value

```
CREATE FUNCTION acct_total(cust_name VARCHAR(20))
RETURNS NUMERIC(12,2)
BEGIN
    -- Variables to accumulate into
    DECLARE bal NUMERIC(12,2);
    DECLARE total NUMERIC(12,2) DEFAULT 0;

    -- Cursor, and flag for when fetching is done
    DECLARE done INT DEFAULT 0;
    DECLARE cur CURSOR FOR
        SELECT balance
        FROM account NATURAL JOIN depositor AS d
        WHERE d.customer_name = cust_name;
```

# Total Account Balance (2)

30

```
-- When fetch is complete, handler sets flag
-- 02000 is MySQL error for "zero rows fetched"
DECLARE CONTINUE HANDLER FOR SQLSTATE '02000'
    SET done = 1;

OPEN cur;
REPEAT
    FETCH cur INTO bal;
    IF NOT done THEN
        SET total = total + bal;
    END IF;
UNTIL done END REPEAT;
CLOSE cur;
RETURN total;

END
```

# Using Our User-Defined Function

31

- Can compute total balances now:

```
SELECT customer_name,  
       acct_total(customer_name) AS total  
FROM customer;
```

□ Result:

customer_name	total
Adams	0.00
Brooks	0.00
Curry	0.00
Glenn	0.00
Green	0.00
Hayes	900.00
Jackson	0.00
Johnson	1400.00
Jones	750.00
Lindsay	700.00
Majeris	850.00
McBride	0.00
Smith	1325.00
Turner	350.00
Williams	0.00

# Stored Routine Benefits

32

- Very effective for manipulating large datasets in unusual ways, within the database
  - ▣ Don't incur communications overhead of sending commands and exchanging data
  - ▣ Database can frequently perform such tasks more efficiently than the applications can
- Often used to provide a secure interface to data
  - ▣ e.g. banks will lock down data tables, and only expose certain operations through stored procedures
- Can encapsulate business logic in procedures
  - ▣ Forbid invalid states by requiring all operations go through stored procedures



# Stored Routine Drawbacks

33

- Increases load on database system
  - ▣ Can reduce performance for *all* operations being performed by DBMS
  - ▣ Need to make sure the operation *really* requires a stored procedure...
    - Most projects do not need stored procedures!
- Very hard to migrate to a different DBMS
  - ▣ Different vendors' procedural languages have *many* distinct features and limitations